

# Logistic Regression

# Classification

Email: Spam / Not Spam?


Online Transactions: Fraudulent (Yes / No)?

Tumor: Malignant / Benign?

$y \in \{0,1\}$

- 0: Negative Class (benign tumor)
- 1: Positive Class (malignant tumor)

→ **Binary classification**



Note:  $y \in \{0,1,2,3, \dots\}$ : **Multi-class classification** is an extension of binary classification


## Classification Based on Probability

Instead of just predicting the class, give the probability of the instance being that class, i.e., learn  $p(y|x)$

Recall that:

$$0 \leq p(\text{event}) \leq 1$$

$$p(\text{event}) + p(\neg \text{event}) = 1$$



## Interpretation of Hypothesis Output

$$h_{\theta}(\mathbf{x}) = \text{estimated } p(y = 1 \mid \mathbf{x}; \theta)$$

Example: Cancer diagnosis from tumor size

$$\mathbf{x} = \begin{bmatrix} x_0 \\ x_1 \end{bmatrix} = \begin{bmatrix} 1 \\ \text{tumorSize} \end{bmatrix}$$

$$h_{\theta}(\mathbf{x}) = 0.7$$

→ Tell patient that 70% chance of tumor being malignant


Note that:  $p(y = 0 \mid \mathbf{x}; \theta) + p(y = 1 \mid \mathbf{x}; \theta) = 1$

Therefore,  $p(y = 0 \mid \mathbf{x}; \theta) = 1 - p(y = 1 \mid \mathbf{x}; \theta)$

# Logistic Regression

$$h_{\theta}(x) = g(\theta^T x)$$


$$g(z) = \frac{1}{1 + e^{-z}}$$



$\theta^T x$  should be large negative values for negative instances

$\theta^T x$  should be large positive values for positive instances

- Assume a threshold and...
  - Predict  $y = 1$  if  $h_{\theta}(x) \geq 0.5$
  - Predict  $y = 0$  if  $h_{\theta}(x) < 0.5$



## Classification

Classification:  $y=0$  or  $y=1$ , but  $h_{\theta}(x)$  can be  $>1$  or  $<0$

Logistic regression:  $0 \leq h_{\theta}(x) \leq 1$


→ use **Sigmoid / Logistic Function**

# Logistic Regression Model

We want our classifier to output values between 0 and 1

- When using linear regression we did  $h_{\theta}(x) = \theta^T x$
- For classification hypothesis representation we do  $h_{\theta}(x) = g(\theta^T x)$  where  $g(z) = \frac{1}{1+e^{-z}}$  is a Sigmoid (or Logistic) function.

$$\text{Thus } h_{\theta}(x) = \frac{1}{1+e^{-\theta^T x}}$$




## Logistic Regression

- Takes a probabilistic approach to learning discriminative functions (i.e., a classifier)
- $h_{\theta}(x)$  should give  $p(y = 1 \mid x; \theta)$ 
  - Want  $0 \leq h_{\theta}(x) \leq 1$
- Logistic regression model:

$$h_{\theta}(x) = g(\theta^T x)$$

$$g(z) = \frac{1}{1 + e^{-z}}$$

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$





# Logistic Regression

Training  
set:

$$\{(x^{(1)}, y^{(1)}), (x^{(2)}, y^{(2)}), \dots, (x^{(m)}, y^{(m)})\}$$

m examples

$$x \in \begin{bmatrix} x_0 \\ x_1 \\ \dots \\ x_n \end{bmatrix} \mathbb{R}^{n+1}$$

$$\underline{x_0 = 1}, \underline{y \in \{0, 1\}}$$

$$\left[ h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}} \right]$$

How to choose parameters  $\theta$  ?



## A **non-convex** loss function with MSE cost function

Linear regression:  $J(\theta) = \frac{1}{m} \sum_{i=1}^m \frac{1}{2} \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$

$$\text{Cost} \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) = \frac{1}{2} \left( h_{\theta}(x^{(i)}) - y^{(i)} \right)^2$$

But if  $h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$ , the cost function is **non-convex** (with MSE).

→ Easy to be trapped at a local minimum → try to make it **convex**.



## A **convex** logistic regression cost function

$$\text{cost}(h_{\boldsymbol{\theta}}(\mathbf{x}), y) = \begin{cases} -\log(h_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\boldsymbol{\theta}}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$


This is the penalty the algorithm pays.

## Intuition Behind the Objective Function

$$\text{cost}(h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$


If  $y = 1$

- Cost = 0 if prediction is correct
- As  $h_{\theta}(\mathbf{x}) \rightarrow 0$ ,  $\text{cost} \rightarrow \infty$
- Captures intuition that larger mistakes should get larger penalties
  - e.g., predict  $h_{\theta}(\mathbf{x}) = 0$ , but  $y = 1$



## Intuition Behind the Objective Function

$$\text{cost}(h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$



If  $y = 0$

- Cost = 0 if prediction is correct
- As  $(1 - h_{\theta}(\mathbf{x})) \rightarrow 0$ ,  $\text{cost} \rightarrow \infty$
- Captures intuition that larger mistakes should get larger penalties

## Cost Function Simplification

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{cost} \left( h_{\theta}(x^{(i)}), y^{(i)} \right)$$

$$\text{cost} (h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

Note:  $y = 0$  or  $y = 1$  always

How to rewrite (simplify) the cost function  $J(\theta)$ ?

## Cost Function Simplification

$$J(\theta) = \frac{1}{m} \sum_{i=1}^m \text{cost} \left( h_{\theta}(x^{(i)}), y^{(i)} \right)$$

$$\text{cost} (h_{\theta}(\mathbf{x}), y) = \begin{cases} -\log(h_{\theta}(\mathbf{x})) & \text{if } y = 1 \\ -\log(1 - h_{\theta}(\mathbf{x})) & \text{if } y = 0 \end{cases}$$

Note:  $y = 0$  or  $y = 1$  always

$$\text{cost} (h_{\theta}(x), y) = -y \log(h_{\theta}(x)) - (1 - y) \log(1 - h_{\theta}(x))$$

$$\text{If } y = 1 : \text{cost} (h_{\theta}(x), y) = -\log(h_{\theta}(x))$$

$$\text{If } y = 0 : \text{cost} (h_{\theta}(x), y) = -\log(1 - h_{\theta}(x))$$

## Logistic Regression Cost Function

$$\begin{aligned} J(\theta) &= \frac{1}{m} \sum_{i=1}^m \text{cost} \left( h_{\theta}(x^{(i)}), y^{(i)} \right) \\ &= -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right] \end{aligned}$$

To fit parameters  $\theta$ :

Compute  $\min_{\theta} J(\theta) \rightarrow$  Get  $\theta$

To make prediction given new  $x$ :

$$\text{Output } h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}} = p(y = 1|x, \theta)$$



## Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right]$$

Want  $\min_{\theta} J(\theta)$ :

Repeat {

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

}

where 
$$\frac{\partial}{\partial \theta_j} J(\theta) = \frac{1}{m} \sum_{i=1}^m \left( h_{\theta}(x^{(i)}) - y^{(i)} \right) x_j^{(i)}$$

## Gradient Descent

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right]$$

Want  $\min_{\theta} J(\theta)$ :

Repeat {

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)}$$

}

This looks IDENTICAL to linear regression!!!

- Ignoring the  $1/m$  constant
- However, the form of the model is very different:

$$h_{\theta}(x) = \frac{1}{1 + e^{-\theta^T x}}$$

## Gradient Descent with Regularization

$$J(\theta) = -\frac{1}{m} \left[ \sum_{i=1}^m y^{(i)} \log(h_{\theta}(x^{(i)})) + (1 - y^{(i)}) \log(1 - h_{\theta}(x^{(i)})) \right]$$

$$J_{regularized}(\theta) = J(\theta) + \lambda \sum_{j=1}^d \theta_j^2 = J(\theta) + \lambda \|\theta_{[1:d]}\|_2^2$$

Want  $\min_{\theta} J(\theta)$ :


Repeat {

$$\theta_j := \theta_j - \alpha \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x_j^{(i)} - \lambda \theta_j$$


}

## Multi-class classification

Binary classification:




Multi-class classification:



Disease diagnosis: healthy / cold / flu / pneumonia

Object classification: desk / chair / monitor / bookcase


## One-vs-all (one-vs rest)



Take the max-probability class among all logistic regression classifiers.  
Extra reading: **softmax regression**.

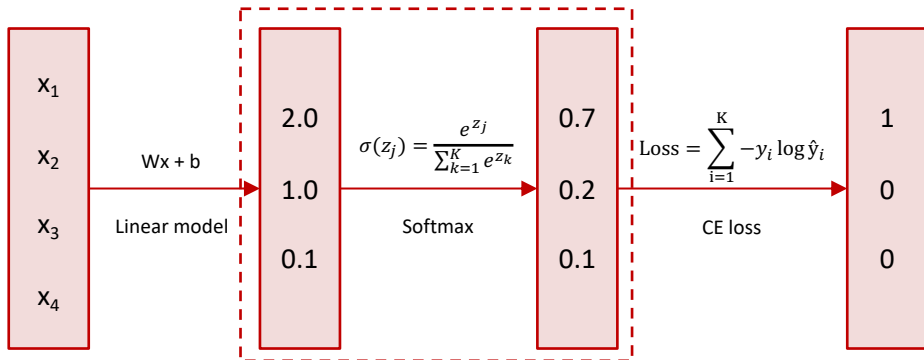
# Softmax regression

**Softmax regression** (or multinomial logistic regression) is a generalization of logistic regression to the case where we want to handle multiple classes.



# Softmax regression

**Softmax regression** (or multinomial logistic regression) is a generalization of logistic regression to the case where we want to handle multiple classes.



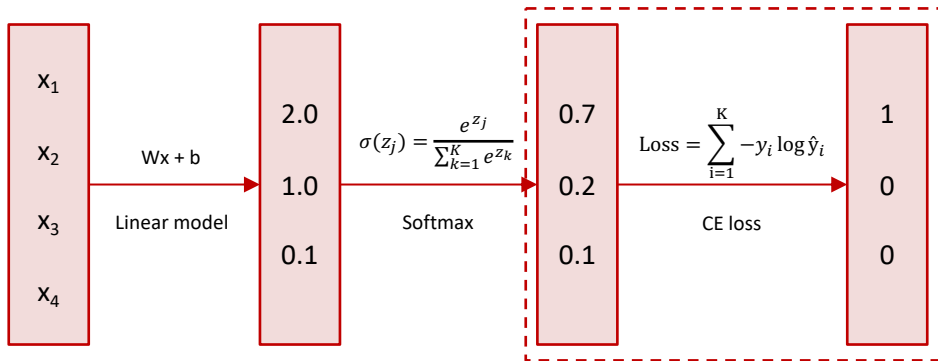
$$\sigma(2.0) = \frac{e^{2.0}}{e^{2.0} + e^{1.0} + e^{0.1}} = 0.7$$

$$\sigma(0.1) = \frac{e^{0.1}}{e^{2.0} + e^{1.0} + e^{0.1}} = 0.1$$

$$\sigma(1.0) = \frac{e^{1.0}}{e^{2.0} + e^{1.0} + e^{0.1}} = 0.2$$

# Softmax regression

**Softmax regression** (or multinomial logistic regression) is a generalization of logistic regression to the case where we want to handle multiple classes.



$$\text{Loss} = -1 \times \log_2 0.7 - 0 \times \log_2 0.2 - 0 \times \log_2 0.1 = 0.51$$



## Evaluation metrics

General method: calculate the **difference** between ground-truth labels and model predictions.

Example: testing 165 emails in a spam/non-spam classification problem.

	Prediction YES	Prediction NO
Actual YES	100	5
Actual NO	10	50

## Evaluation metrics

Example: testing 165 emails in a spam/non-spam classification problem.

	<b>Prediction YES</b>	<b>Prediction NO</b>
<b>Actual YES</b>	100	5
<b>Actual NO</b>	10	50

- Precision =  $100/(100+10) \sim 91\%$ : how many predicted items are relevant.
- Recall =  $100/(100+5) \sim 95\%$ : how many relevant items are predicted.

## Evaluation metrics

Example: testing 165 emails in a spam/non-spam classification problem.

	<b>Prediction YES</b>	<b>Prediction NO</b>
<b>Actual YES</b>	True Positive TP	False Negative FN
<b>Actual NO</b>	False Positive FP	True Negative TN

$$Precision = \frac{TP}{TP + FP}$$

$$Recall = \frac{TP}{TP + FN}$$

$$F1 = \frac{2 \times Precision \times Recall}{Precision + Recall}$$

## Summary

Binary Classification

Decision Boundary

Logistic Regression

- Sigmoid function
- Cost Function
- Optimization
- Regularization

Multi-class (Multinomial Classification)

- One-vs-all
- Softmax regression

Evaluation metrics

Thank you